

# Introduction to Machine Learning

---

**Inteligencia Artificial**

**3° INF**



# Learning agents

---

## ■ Learning from experience

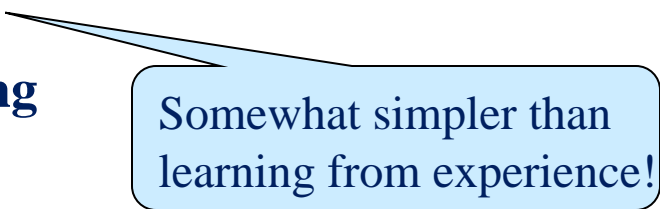
Interaction with the surroundings + other agents

- **Learning by imitation** (needs interaction with instructor).
- **Reinforcement learning:** Trial & error + reward/punishment.

## ■ Learning from data

### ■ Unsupervised learning

- **Clustering**
- **Estimation of densities (e.g.. Generative models)**
- *Autoencoders*
- **Anomaly detection**



Somewhat simpler than  
learning from experience!

### ■ Supervised learning

### ■ Semi-supervised learning

# Reinforcement learning

---

- **Goal:** Learn how to make sequence of decisions.
- **Method:** Trial and error
  - The agent carries out **sequence of actions**  
Example: moves in a chess game
  - **At the end of each sequence** agent receives a **reward** or a **penalty**.  
Example: Reward in agent wins match  
Penalty if agent loses
  - **Credit assignment problem:** How does one determine impact of each of the decisions in the sequence in the final result?  
**Solution:** Compute a “**value function**” at each state that estimates the optimal (final) reward from that state.

# AlphaGo Zero

---

- **Initially** the agent **has no knowledge** of **Go**, except for its rules.
- The agent is equipped with a
  - **Advanced search tree**
  - Two **neuronal networks**, which take as input a description of its board at the current state.
    - The **policy network** selects the next move.
    - The **value network** predicts the winner of the game from that state.
- The **reinforcement learning process** consists in **modifying the weights** of the networks **using the reward / penalty** that corresponds to winning or losing a game.
- The **system** progressively **improves** the quality of its moves by **playing Go with modified copies** of itself.

<https://deepmind.com/blog/alphago-zero-learning-scratch/#gif-120>

# Inductive learning from data

---

- 1. Data collection:** Capture, encoding and storage of instances (= examples) from which one wishes to learn.
- 2. Feature selection and feature construction:** For each of the instances, one determines the features that characterized de data.
- 3. Choice of learning algorithm:**
  1. Unsupervised: K-means, fuzzy K-means, spectral clustering, etc.
  2. Supervised: Naïve Bayes, nearest-neighbors, decision trees, ensembles (*bagging, boosting, random forest*), support vector machines, neural networks, Gaussian processes, etc.
- 4. Training:** Application of a learning algorithm on the labeled data to determine the parameters of the system that performs clustering or prediction (typically, an optimization process).
- 5. Validation:** Process by which the quality of the system that has been trained is assessed.

# Learning by induction from data

---

- **Unsupervised:** Inference from **unlabeled data**  $\{\mathbf{x}_n\}_{n=1}^N$ 
  - **Clustering** (hierarchical, K-means, fuzzy K-means, etc.)
  - **Probability modeling** (density estimation)
  - **Autoencoders**
  - **Anomaly detection**
- **Supervised:** Learning to predict from a **set of labeled instances** (aka examples, instances, samples):  $\{(\mathbf{x}_n, y_n)\}_{n=1}^N$ 
  - $\mathbf{x} \in \mathcal{X}$ : vector of **attributes**, features, independent variables, input variables. covariates...
  - $y \in \mathcal{Y}$ : (class) **label**, dependent variable, outcome, target, ...
- **Semisupervised:** Learning to predict from data that are **partially labeled**.

# Unsupervised learning

---

Given a set of unlabeled data

$$\mathcal{D}_{train} = \{(\mathbf{x}_n)\}_{n=1}^{N_{train}};$$

$\mathbf{x}_n = (x_{n1} \dots x_{nD})^T \in \mathcal{X}$  : Vector of attributes  
(inputs, features, independent variables, etc.).

- **Clustering:** Identification of natural groupings
  - *K-means*
  - *Fuzzy K-means*
  - *Hierarchical clustering*
- **Density estimation / generative models**
  - Mixture models (e.g. Gaussian mixtures)
- **Anomaly detection**

# Clustering for image segmentation



**Figure 9.3** Two examples of the application of the  $K$ -means clustering algorithm to image segmentation showing the initial images together with their  $K$ -means segmentations obtained using various values of  $K$ . This also illustrates the use of vector quantization for data compression, in which smaller values of  $K$  give higher compression at the expense of poorer image quality.

Image source: Christopher M. Bishop. 2006. Pattern Recognition and Machine Learning (Information Science and Statistics). Springer-Verlag, Berlin, Heidelberg.



# Labeled dataset

---

Learning to predict from labeled data  $\mathcal{D} = \{(\mathbf{x}_n, y_n)\}_{n=1}^N$

$\mathbf{x}_n^T = (x_{n1} \ x_{n2} \ \dots \ x_{nD})$  [ $D$ -dimensional vector of attributes]

- **Regression:** **Labels**  $y$  are **continuous** (e. g.  $y \in \mathbb{R}$ )
- **Classification:** **Labels**  $y$  are **discrete (no ordering)**  
 $y \in \{C_1, C_2, \dots, C_K\}$
- **Ordinal classification:** **Labels**  $y$  are **discrete & ordered**  
 $y \in \{C_1, C_2, \dots, C_K\};$   
 $C_1 \leq C_2 \leq \dots \leq C_K$

# Supervised learning

---

Given a labeled dataset:  $\mathcal{D}_{train} = \{(\mathbf{x}_n, y_n)\}_{n=1}^{N_{train}}$ ;

$\mathbf{x}_n = (x_{n1} \dots x_{nD})^T \in \mathcal{X}$  : Vector of attributes  
(inputs, characteristics, independent variables,...)

$y_n \in \mathcal{Y}$ : Variable to predict

(outputs, label, dependent variable,...)

- **Objective:** **Induce** from  $\mathcal{D}_{train}$  **a function** to **predict** the value of  $y$  from the value of the vector of attributes  $\mathbf{x}$ . Predictions should be accurate for instances that could be different from the training ones (**generalization capacity**).

- **Learning algorithm**

$\mathcal{L}: \mathcal{D}_{train} \rightarrow h$  [predictor (model) induced from  $\mathcal{D}_{train}$ ]

$h: \mathbf{x} \in \mathcal{X} \rightarrow h(\mathbf{x}) \in \mathcal{Y}$

# Data preparation: Attributes

---

The vector attributes (or features)  $\mathbf{x}$  characterizes the instances whose class label we want to predict.

- **Feature construction**
- **Feature (variable) selection**
- **Types of attributes**
  - **Ordinal:** discrete / continuous
  - **Nominal:** There is no order relation between the values this type of attribute can take

1-of-K encoding

$$C_1 \equiv (1 \ 0 \ \dots \ 0)$$

$$C_2 \equiv (0 \ 1 \ \dots \ 0)$$

...

$$C_K \equiv (0 \ 0 \ \dots \ 1)$$

Needed for SVMs,  
neural networks, etc.  
Not needed for  
decision trees

# Data preparation: Preprocessing

Training data:  $\mathcal{D}_{train} = \{(\mathbf{x}_n, y_n)\}_{n=1}^{N_{train}}$

■ **Outlier detection**

■ **Handling of missing values**

■ **Centering and scaling**  $x_{nd} \rightarrow \frac{x_{nd} - m_d}{s_d}$

IMPORTANT:  
Estimates in  $\mathcal{D}_{train}$   
(no peeking at test)

Center: $m_d$	Scale: $s_d$
Mean: $\bar{x}_d = \frac{1}{N_{train}} \sum_{n=1}^{N_{train}} x_{nd}$	stdev = $\sqrt{\frac{1}{N_{train}} \sum_{n=1}^{N_{train}} (x_{nd} - \bar{x}_d)^2}$
Median	Interquartile range
$\frac{\max\{x_{nd}\}_{n=1}^{N_{train}} + \min\{x_{nd}\}_{n=1}^{N_{train}}}{2}$	$\frac{\max\{x_{nd}\}_{n=1}^{N_{train}} - \min\{x_{nd}\}_{n=1}^{N_{train}}}{2}$

# Preprocessing in scikit-learn

---

<http://scikit-learn.org/stable/modules/preprocessing.html>

- StandardScaler:      Centering with the mean  
                                Scaling with the standard deviation
- RobustScaler:        Centering with the median  
                                Scaling with the interquartile range
- MaxAbsScaler:        Data normalized in  $[-1,1]$  (does not center data)
- MinMaxScaler:        Data normalized in  $[0,1]$

Should I normalize/standardize/rescale the data?

<http://www.faqs.org/faqs/ai-faq/neural-nets/part2/section-16.html>

# Data preparation: Class label encoding

Learning to predict from labeled data  $\{(\mathbf{x}_n, c_n)\}_{n=1}^N$

■ **Binary classification:**  $c_n \in \{C_0, C_1\}$  Encoding  $\{(\mathbf{x}_n, t_n)\}_{n=1}^N$

- **Zero-one encoding:**  $t_n = \mathbb{I}[c_n = C_1] = \begin{cases} 0 & \text{if } c_n = C_0 \\ 1 & \text{if } c_n = C_1 \end{cases}$
- **Sign encoding:**  $t_n = 2\mathbb{I}[c_n = C_1] - 1 = \begin{cases} -1 & \text{if } c_n = C_0 \\ 1 & \text{if } c_n = C_1 \end{cases}$

■ **Multiclass classification:**  $y_n \in \{C_1, C_2, \dots, C_K\}$

$$\{(\mathbf{x}_n, \mathbf{t}_n)\}_{n=1}^N; \quad \mathbf{t}_n^T = (t_{n1} \ t_{n2} \ \dots \ t_{nK})$$

$$t_{nk} = \mathbb{I}[c_n = C_k] = \begin{cases} 0 & \text{if } c_n \neq C_k \\ 1 & \text{if } c_n = C_k \end{cases}; \quad k = 1, 2, \dots, K$$

E.g.  $K = 3$      $C_1: 100, C_2: 010, C_3: 001$  [**1-of-K or 1-hot encoding**]

Vertices of a regular (probability) simplex.

# $k$ -Nearest neighbors ( $k$ -NN)

## ■ Input:

Vector of attributes:  $\mathbf{x}_n \in \mathcal{X}$

Class label:  $c_n \in \mathcal{C}$

■ Training data:  $\mathcal{D}_{train} = \{(\mathbf{x}_n, c_n)\}_{n=1}^{N_{train}}$

■ Test instance:  $\mathbf{x}_{test}$

■ Nr. of neighbors:  $k > 0$  (typically uneven, to avoid ties)

■ Metric to compute distances:  $d(\mathbf{x}_i, \mathbf{x}_j)$

E.g. Euclidean distance,  
Manhattan distance, etc.

■ Output:  $h_{kNN}(\mathbf{x}_{test}; \mathcal{D}_{train})$

## ■ Pseudocode:

Class label prediction for  $\mathbf{x}_{test}$ .


1. Find the  $k$  **nearest-neighbors** to  $\mathbf{x}_{test}$  **in the training set**, using the **distance function**  $d(\mathbf{x}_i, \mathbf{x}_{test})$ .

2. **Assign to  $\mathbf{x}_{test}$  the majority class** among the  $k$  neighbors.

■ Note: Ties can be resolved in a number of ways. For instance, by reducing the value of  $k$  by one unit at a time until the tie is resolved.

# Example: Iris dataset ( $k = 1$ )

- Attributes: “Petal length”, “Sepal length”
- Classes: “Setosa”, “Versicolor”, “Virginica”



In the original Iris problem, the vector of attributes has more components. For the purpose of visualization, we consider only two in this simplified classification problem

- $k = 1$  Assign “virginica”, which is the class of the nearest neighbor to  $\mathbf{x}_{test}$

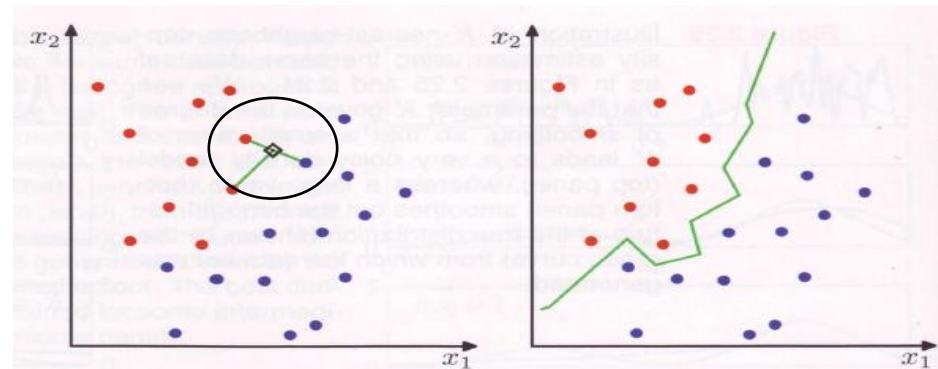


# $k$ -NN and Bayes theorem

Let us consider the smallest ball around  $\mathbf{x}_{test}$  that contains exactly  $k$  instances from the training set

$k = 3$

[Fig. 2.27, Bishop]



$k$ -NN can be seen as a MAP (maximum a posteriori) predictor that uses **local estimations** of the relevant probability densities

$$P(c|\mathbf{x}_{test}) \approx \frac{N_c(\mathbf{x}_{test})}{N(\mathbf{x}_{test})}$$

$k$ -NN prediction

number of instances in the ball centered in  $\mathbf{x}_{test}$

number of instances of class  $c$  in the same ball

$$c_{MAP}^*(\mathbf{x}_{test}) = \arg \max_{c \in \mathcal{C}} P(c|\mathbf{x}_{test}) = \arg \max_{c \in \mathcal{C}} N_c(\mathbf{x}_{test})$$

# How many neighbors in $k$ -NN?

The number of neighbors can be seen as a smoothing parameter: The larger  $k$  is, the smoother the classification boundary

[Fig. 2.16, James et al.]

- Asymptotic properties  $N_{train} \rightarrow \infty$

$$k = 1: \text{error}_{1\text{-NN}} < 2 \text{error}_{\text{Bayes}}$$

$$k = 3: \text{error}_{3\text{-NN}} \approx \text{error}_{\text{Bayes}} + 3 (\text{error}_{\text{Bayes}})^2$$

- Estimation of  $k$  by **leave-one-out (loo) cross-validation**

$$\text{err}_{k\text{NN}}(\mathbf{x}_n; k, \mathcal{D}_{train \setminus n}) = \mathbb{I}[c_n \neq h_{k\text{NN}}(\mathbf{x}_n; \mathcal{D}_{train \setminus n})]$$

$$k_{loo}^* = \arg \min_k \frac{1}{N_{train}} \sum_{n=1}^{N_{train}} \text{err}_{k\text{NN}}(\mathbf{x}_n; k, \mathcal{D}_{train \setminus n})$$

Leave-one-out error estimate

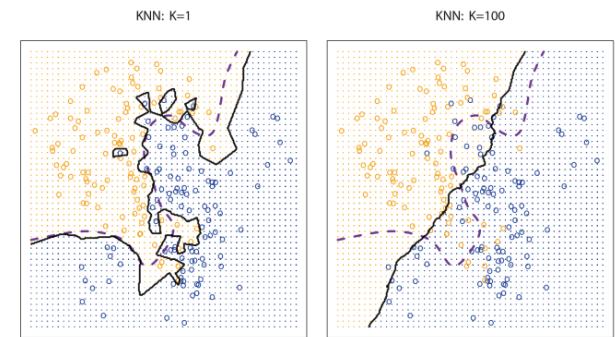


FIGURE 2.16. A comparison of the KNN decision boundaries (solid black curves) obtained using  $K = 1$  and  $K = 100$  on the data from Figure 2.13. With  $K = 1$ , the decision boundary is overly flexible, while with  $K = 100$  it is not sufficiently flexible. The Bayes decision boundary is shown as a purple dashed line.

$$\mathcal{D}_{train \setminus n} = \{(\mathbf{x}_m, c_m)\}_{m=1, m \neq n}^{N_{train}}$$

# k-NN: pros & cons

---

## ■ Advantages

- Simple implementation.
- Typically, good results if appropriate distance function is chosen.
- It is possible to provide an interpretation of the prediction results in terms of the examples used to determine the class.

## ■ Drawbacks

- Difficulties with irrelevant or noisy attributes.
- Selection of  $k$ .
- Selection of distance function
  - Scaling of attributes
  - Handling of nominal attributes.
- High computational cost at prediction time.