

Decision trees

Artificial Intelligence

3rd year INF

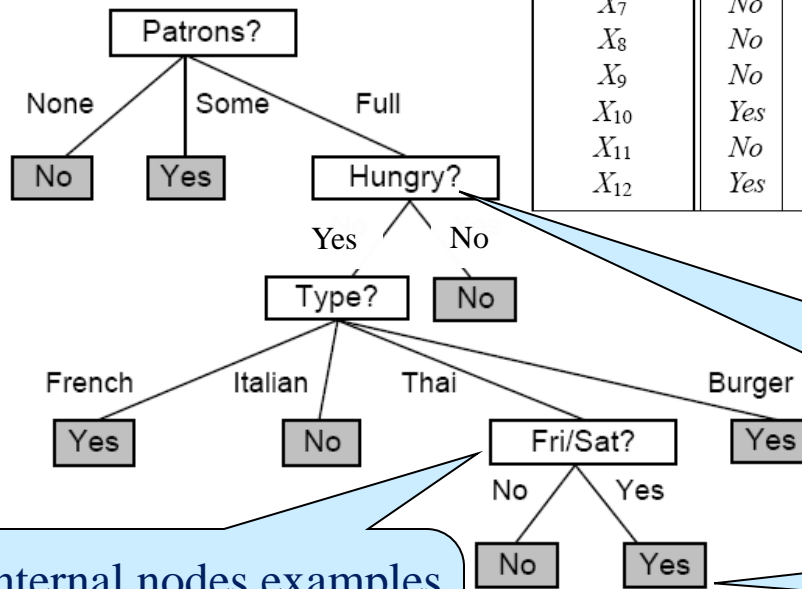


Decision tree

Training data

Learned tree:

Example	Attributes										Goal
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>WillWait</i>
X_1	Yes	No	No	Yes	Some	\$\$\$	No	Yes	French	0–10	Yes
X_2	Yes	No	No	Yes	Full	\$	No	No	Thai	30–60	No
X_3	No	Yes	No	No	Some	\$	No	No	Burger	0–10	Yes
X_4	Yes	No	Yes	Yes	Full	\$	No	No	Thai	10–30	Yes
X_5	Yes	No	Yes	No	Full	\$\$\$	No	Yes	French	>60	No
X_6	No	Yes	No	Yes	Some	\$\$	Yes	Yes	Italian	0–10	Yes
X_7	No	Yes	No	No	None	\$	Yes	No	Burger	0–10	No
X_8	No	No	No	Yes	Some	\$\$	Yes	Yes	Thai	0–10	Yes
X_9	No	Yes	Yes	No	Full	\$	Yes	No	Burger	>60	No
X_{10}	Yes	Yes	Yes	Yes	Full	\$\$\$	No	Yes	Italian	10–30	No
X_{11}	No	No	No	No	None	\$	No	No	Thai	0–10	No
X_{12}	Yes	Yes	Yes	Yes	Full	\$	No	No	Burger	30–60	Yes



Instances for which
Patrons = full
are associated to this node

At internal nodes examples
are separated according to
a test on an attribute

Instances that reach this
leaf node are classified as
WillWait = Yes

How does it work?

- A decision tree is a **hierarchical questionnaire** that **splits the data** according to a **sequence of tests on their attributes**.
- **Each example**, when processed by the tree, follows a unique **path** from the **root node** to the corresponding **leaf** according to the results of the **tests on the attributes** performed at each of the **intermediate internal nodes**.
- The **class associated to a node** corresponds to the **majority label of the training instances** assigned to that node.
- The **tests** at the internal nodes are determined by **maximizing a quantity** (e.g. the information gain) that favors a **clearer separation of the classes** in the children of such nodes.
- The **class label prediction** for an example takes place at the corresponding **leaf node**.

Learning algorithm

function DECISION-TREE-LEARNING(*examples*, *attributes*, *default*) **returns** a decision tree

inputs: *examples*, set of examples
attributes, set of attributes
default, default value for the goal predicate

Simplified version of
Quinlan's ID3 (1986)

if *examples* is empty **then return** *default*

else if all *examples* have the same classification **then return** the classification

else if *attributes* is empty **then return** MAJORITY-VALUE(*examples*)

else

best \leftarrow CHOOSE-ATTRIBUTE(*attributes*, *examples*)

tree \leftarrow a new decision tree with root test *best*

for each value v_i of *best* **do**

examples_i \leftarrow {elements of *examples* with *best* = v_i }

subtree \leftarrow DECISION-TREE-LEARNING(*examples_i*, *attributes* – *best*,
MAJORITY-VALUE(*examples*))

add a branch to *tree* with label v_i and subtree *subtree*

end

return *tree*

The **best attribute** is the
one that provides the largest
amount of **information** on
the class label.

Recursion

Measuring information: Binary Entropy

“r.v.” means
“random variable”

“w.p.” means
“with probability”

Consider the binary r.v. $X = \begin{cases} x_1 & \text{w.p. } p \\ x_0 & \text{w.p. } q = 1 - p \end{cases}$


Entropy
of r.v. X

$$H(X) = H_b(p) = -p \log_2 p - q \log_2 q \text{ [in bits]}$$

$$p = P[X = x_1] \\ q = P[X = x_0]$$

$$0 \leq p, q \leq 1 \\ p + q = 1$$

$0 \leq H(X) \leq 1$
Information cannot
be negative.



A common
unit to measure
information

Another, less common
measurement unit
for information

Also: $H(X) = -p \log p - q \log q$ [in nats]

Natural logarithm

Message encoding: Entropy rate

$x_1 = 'a'$
 $x_0 = 'b'$

■ Information contents of messages with random $\{ 'a', 'b' \}$:

■ Message 1: “aaaaaaaaaaaaaaaaaaaaa”

$H_b(p)$ is minimum
for $p = 1, q = 0$
 $p = 0, q = 1$

$$\hat{p} = 1; \hat{q} = 0$$

Estimate the probabilities
from the frequencies of the
symbols in the message

$$H(X) = 0 \text{ bits}$$

Send message: “20 a’s”
When length of the actual message
is very large, the average amount of
information per symbol that needs
to be transmitted approaches 0 bits

■ Message 2: “aaababaaabaaabaaaaab”

$$\hat{p} = \frac{3}{4}; \hat{q} = \frac{1}{4}$$

$$H(X) = 0.81 \text{ bits}$$

■ Message 3: “abbababbaababbaaabba”

$H_b(p)$ is maximum
for $p = q = \frac{1}{2}$

$$\hat{p} = \frac{1}{2}; \hat{q} = \frac{1}{2}$$

$$H(X) = 1 \text{ bit}$$

All symbols need to be transmitted:
On average 1 bit per symbol

Entropy for a discrete r.v.

■ Consider the discrete r.v.

\mathcal{X} is the space in which the random variable takes values

$$X \in \mathcal{X} = \{x_1, x_2, \dots, x_{|\mathcal{X}|}\}$$

Maximum: $H(X) = \log_2 |\mathcal{X}|$
for $P[X = x_i] = \frac{1}{|\mathcal{X}|}$, $i = 1, 2, \dots, |\mathcal{X}|$

Random variable X can take $|\mathcal{X}|$ different values

$$H(X) = - \sum_{i=1}^{|\mathcal{X}|} P[X = x_i] \log_2 P[X = x_i] \quad [\text{in bits}]$$

If a sender wants to transmit values (sampled independently) of the r.v. X to a receiver, the entropy measures the average amount of bits per symbol of the minimal length message

Message encoding: Entropy rate

$x_1 = 'a'$
 $x_2 = 'b'$
 $x_3 = 'c'$
 $x_4 = 'd'$

■ Information contents of messages with {'a','b','c','d'}:

■ **Message 1:** $p('a') = p('b') = p('c') = p('d') = 1/2$

$H(X) = 2 \text{ bits}$

Symbol	'a'	'b'	'c'	'd'
Encoding	00	01	10	11

Fixed length code

■ **Message 2:** $p('a') = \frac{1}{2}$; $p('b') = p('c') = \frac{1}{4}$; $p('d') = 0$

$H(X) = 1.5 \text{ bits}$

Symbol	'a'	'b'	'c'	'd'
Encoding	0	10	11	-

Variable length
prefix code

• Average # bits per symbol of encoded message:

$$p('a') \times 1 \text{ bit} + p('b') \times 2 \text{ bit} + p('c') \times 2 \text{ bits} + p('d') \times 0 \text{ bits} = 1.5 \text{ bits}$$

An alternative: Gini impurity


- Consider the discrete r.v. $X \in \mathcal{X} = \{x_1, x_2, \dots, x_{|\mathcal{X}|}\}$

$$I_G(X) = 1 - \sum_{i=1}^{|\mathcal{X}|} (P[X = x_i])^2$$

Used to determine splits
in CART decision
trees (Breiman et al. 1984)

- Gini Impurity for a binary r.v. $I_G(X) = 1 - p^2 - q^2$

Similar shape to the
binary entropy



Conditional entropy

- **Consider the discrete r.v.'s:**
 $X \in \mathcal{X} = \{x_1, x_2, \dots, x_{|\mathcal{X}|}\}$
 $Y \in \mathcal{Y} = \{y_1, y_2, \dots, y_{|\mathcal{Y}|}\}$

- Entropy of r.v. Y conditioned on $X = x_i$

$$H(Y | X = x_i) = \sum_{j=1}^{|\mathcal{Y}|} P[Y = y_j | X = x_i] \log_2 P[Y = y_j | X = x_i]$$

- **Conditional entropy:**

$$H(Y|X) \leq H(Y)$$

$H(Y|X) = H(Y)$
iff X and Y
are independent

$$H(Y | X) = - \sum_{i=1}^{|\mathcal{Y}|} P[X = x_i] H(Y | X = x_i)$$

Average over the possible values of X

If a sender wants to transmit values of Y , the conditional entropy measures the average number of bits per symbol of the minimal message, assuming the value of X is known by the receiver

Information gain

$$H(Y) \geq H(Y|X)$$

Therefore, $IG(Y | X) \geq 0$

$IG(Y | X) = 0$ iff X and Y are independent

$$IG(Y | X) = H(Y) - H(Y|X) \text{ [in bits]}$$

Measures the average number of bits per symbol of the minimal message to transmit values of the random variable Y that one saves by assuming that the receiver knows the value of X

- Select ***best* attribute** as the one that **maximizes the information gain** of the class given by that attribute.

Split at the root of the tree

How do I determine the first question in the decision tree?

- Class variable: $WillWait \in \{yes, no\}$

Number of instances: $N = 12$ ($N_{yes} = 6$; $N_{no} = 6$)

$$H(WillWait) = H_b\left(\frac{6}{12}\right) = 1 \text{ bit}$$

- The best attribute to make a split at the root of the node is the one that maximizes the Information Gain.

- $IG(WillWait|Type) = 0$ bits

No information is gained

- $IG(WillWait|Patrons) = 0.64$ bits

Largest value of the information gain.
Best attribute: *Patrons*

- ...

Check the other values!

IG of *WillWait* from *Patrons*?

- Attribute: *Patrons* $\in \{none, some, full\}$

- $N_{none} = 2$ ($N_{yes,none} = 0$; $N_{no,none} = 2$)

$$H(WillWait|none) = H_b\left(\frac{0}{2}\right) = 0 \text{ bits}$$

- $N_{some} = 4$ ($N_{yes,some} = 4$; $N_{no,some} = 0$)

$$H(WillWait|some) = H_b\left(\frac{4}{4}\right) = 0 \text{ bits}$$

- $N_{full} = 6$ ($N_{yes,full} = 2$; $N_{no,full} = 4$)

$$H(WillWait|full) = H_b\left(\frac{2}{6}\right) = 0.92 \text{ bits}$$

$$H(WillWait|Patrons) = \frac{2}{12} 0 + \frac{2}{12} 0 + \frac{6}{12} 0.92 = 0.46 \text{ bits}$$

$$IG(WillWait|Patrons) = 1 - 0.46 = 0.64 \text{ bits}$$

Recursion: Split at the node *Patrons = full*

- Training instances at node *Patrons = full*:

$$\{X_2, X_4, X_5, X_9, X_{10}, X_{12}\}$$

$$N = 6 \ (N_{yes} = 2; N_{no} = 4) \Rightarrow$$

$$H(WillWait) = H_b\left(\frac{2}{6}\right) = 0.92 \text{ bits}$$

How do I determine the next question in the decision tree?

Check this!

- The best attribute to make a split at this node is *Hungry*

$$H(WillWait|Hungry) = \frac{4}{6} H_b\left(\frac{2}{4}\right) + \frac{2}{6} H_b\left(\frac{0}{2}\right) = 0.67 \text{ bits}$$

$$IG(WillWait|Hungry) = 0.92 - 0.67 = 0.25 \text{ bits}$$

When does one stop splitting a node?

- The training examples assigned to that node belong to the same class. [The leaf node assigns that class label]
- Node has no examples associated to it. [The leaf node assigns the default class label]
- No more attributes left for splitting the data. [The leaf node assigns the majority class label in that node]
- Prepruning (limit the tree size to avoid **overfitting**)
 - The number of training examples associated to the node is below a threshold.
 - The Impurity Gain is below a threshold.[The leaf node assigns the majority class label in that node]

E.g. Threshold =
 I_G of a random split

Underfitting / overfitting

■ Underfitting


The type of predictor considered **has low expressive capacity**.
In consequence, it is no able to capture the dependencies between the attributes and the variable to be predicted.
The error of the predictor is too high.

■ Overfitting

The type of predictor considered is **too flexible** and learns spurious patterns that are not relevant for prediction (e.g. sampling fluctuations, noise, outliers, etc.).

Training estimate of the expected loss is too optimistic and underestimates the actual error.

Underfitting / overfitting



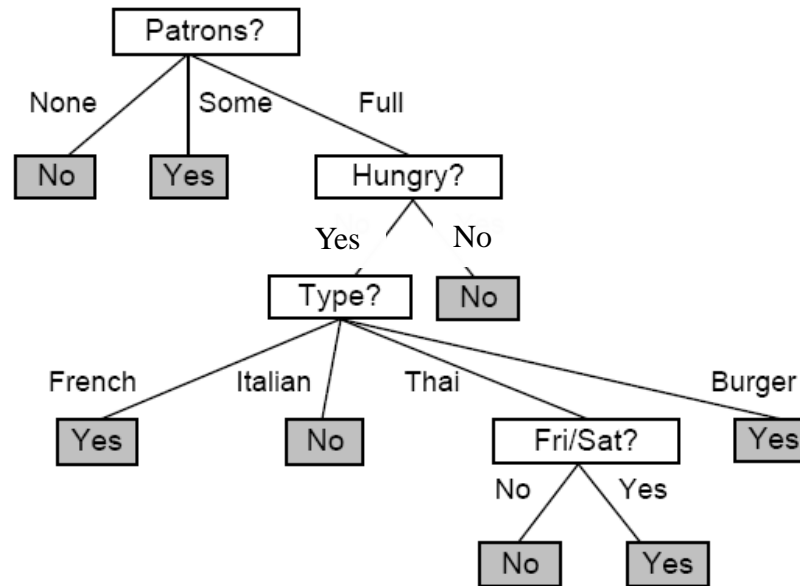
Source: <https://gerardnico.com/> under
license [CC Attribution-Noncommercial-Share Alike 4.0 International](#)

Pruning to avoid overfitting in DT's

- Bias towards smaller (less complex) trees.
 - Prepruning
 - Postpruning: Grow tree to a large size and then prune subtrees that do not provide significant gains in predictive accuracy.
 - Consider an internal node.
 - If turning that node into a leaf does not lead to a significant decrease in the predictive accuracy of the pruned decision tree, then eliminate the subtree which has that node as its root.
 - For this process, accuracy can be estimated on a separate validation set (reduced error pruning), or by CV (e.g. as in CART)
 - Continue pruning until significant deterioration of accuracy

Postpruning is generally preferred. This is a common strategy in machine learning: consider first a potentially complex model and then penalize complexity

Interpretability: Rule extraction



System of rules:

The **group stays** if

either the restaurant has **some patrons**

or (the restaurant is **full** and the **group is hungry** and

(the type of food is **French** or (**Thai** and it is **Fri/Sat**) or **Burger**)

Otherwise, the **group leaves**.


The model is interpretable!

C4.5 Decision tree


- Evolution of ID3 by Quinlan (1992)
- Includes
 - Tests based on numerical attributes
 - Fuzzy decisions
 - Post-pruning
 - Normalization of information gain for multivalued attributes.
 - Handling of missing values
 - Rule extraction and pruning

C4.5: Handling of numerical attributes

- Attributes: N_1, N_2
- Class: “+”, “-” $H(Class) = H_b\left(\frac{8}{16}\right) = 1 \text{ bit}$




Tests on N_1



Pregunta	Rama "No"	Rama "Sí"	Entropía clase en Rama "No"	Entropía clase en Rama "Sí"	H(clase Pre- gunta)	IG
$N_1 > 1$	2+, 1-	6+, 7-	$H(2/3, 1/3) = 0.918$ bits	$H(6/13, 7/13) = 0.996$ bits	$3/16 * 0.918 + 13/16 * 0.996 = 0.981$ bits	$1 - 0.981 = 0.019$ bits
$N_1 > 2$	3+, 2-	5+, 6-	$H(3/5, 2/5) = 0.971$ bits	$H(5/11, 6/11) = 0.994$ bits	$5/16 * 0.971 + 11/16 * 0.994 = 0.987$ bits	$1 - 0.987 = 0.013$ bits
$N_1 > 3$	3+, 4-	5+, 4-	$H(3/7, 4/7) = 0.985$ bits	$H(5/9, 4/9) = 0.991$ bits	$7/16 * 0.985 + 9/16 * 0.991 = 0.988$ bits	$1 - 0.988 = 0.012$ bits
$N_1 > 4$	3+, 5-	5+, 3-	$H(3/8, 5/8) = 0.954$ bits	$H(5/8, 3/8) = 0.954$ bits	$8/16 * 0.954 + 8/16 * 0.954 = 0.954$ bits	$1 - 0.954 = 0.046$ bits
$N_1 > 5$	3+, 6-	5+, 2-	$H(3/9, 6/9) = 0.918$ bits	$H(5/7, 2/7) = 0.863$ bits	$9/16 * 0.918 + 7/16 * 0.863 = 0.894$ bits	$1 - 0.894 = 0.106$ bits
$N_1 > 6$	4+, 7-	4+, 1-	$H(4/11, 7/11) = 0.946$ bits	$H(4/5, 1/5) = 0.722$ bits	$11/16 * 0.946 + 5/16 * 0.722 = 0.876$ bits	$1 - 0.876 = 0.124$ bits
$N_1 > 7$	6+, 7-	2+, 1-	$H(6/13, 7/13) = 0.996$ bits	$H(2/3, 1/3) = 0.918$ bits	$13/16 * 0.996 + 3/16 * 0.918 = 0.981$ bits	$1 - 0.981 = 0.019$ bits
$N_1 > 8$	8+, 8-	0+, 0-	1 bit	--	$16/16 * 1 + 0/16 * -- = 1$ bit	$1 - 1 = 0$ bits



Tests on N_2



Pregunta	Rama "No"	Rama "Sí"	Entropía clase en Rama "No"	Entropía clase en Rama "Sí"	H(clase Pre- gunta)	IG
$N_2 > 21$	2+, 2-	6+, 6-	$H(2/4, 2/4) = 1$ bits	$H(6/12, 6/12) = 1$ bits	$4/16 * 1 + 12/16 * 1 = 1$ bits	$1 - 1 = 0$ bits
$N_2 > 22$	3+, 3-	5+, 5-	$H(3/6, 3/6) = 1$ bits	$H(5/10, 5/10) = 1$ bits	$6/16 * 1 + 10/16 * 1 = 1$ bits	$1 - 1 = 0$ bits
$N_2 > 23$	3+, 4-	5+, 4-	$H(3/7, 4/7) = 0.985$ bits	$H(5/9, 4/9) = 0.991$ bits	$7/16 * 0.985 + 9/16 * 0.991 = 0.988$ bits	$1 - 0.988 = 0.012$ bits
$N_2 > 24$	4+, 6-	4+, 2-	$H(4/10, 6/10) = 0.971$ bits	$H(4/6, 2/6) = 0.918$ bits	$10/16 * 0.971 + 6/16 * 0.918 = 0.951$ bits	$1 - 0.951 = 0.049$ bits
$N_2 > 25$	7+, 8-	1+, 0-	$H(7/15, 8/15) = 0.997$ bits	$H(1/1, 0/1) = 0$ bits	$15/16 * 0.997 + 1/16 * 0 = 0.935$ bits	$1 - 0.894 = 0.065$ bits


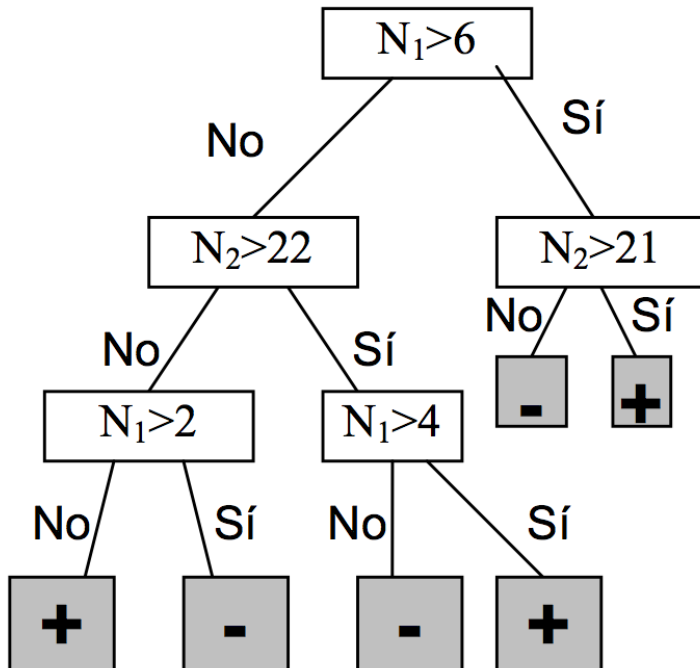
Smaller than
 I_G of test ($N_1 > 6?$)

Test at root node



- The original attribute space has been partitioned into 2 disjoint subspaces (A and B)
- Using a “**divide and conquer**” strategy, and recursively partition A and B separately.

Final C4.5 decision tree



Decision trees: pros & cons

■ Advantages

- Simple implementation.
- Interpretable results.
- Fast training & prediction.

■ Drawbacks

- Not very accurate predictions. However, they can be used as base learners for an ensemble.

<https://scikit-learn.org/stable/modules/tree.html>

Decision forests: Ensembles of DT's

- **Randomization**

- **Bagging**

Bagging and boosting ensembles can also be composed of other types of base learners, such as neural networks

- **Random forest**

- **Randomization + optimization**

- **Boosting**

- **Gradient boosting**

Random forest, gradient boosting, and xgboost have excellent off-the-shelf performance in non-structured problems

- **Xgboost (Extreme Gradient Boosting)**

[<https://xgboost.readthedocs.io/en/latest/>]

[<https://scikit-learn.org/stable/modules/ensemble.html>]